

# Shivam Taneja

Software Engineer

905-914-2526 | [Tanejashivam.st@gmail.com](mailto:Tanejashivam.st@gmail.com) | [Tanejs4.github.io](https://Tanejs4.github.io) | [Linkedin](#)

## Skills

**Cloud Platforms:** AWS (S3, EKS, Lambda), Azure, OpenShift, Terraform (IaC), Docker, Kubernetes

**DevOps & Observability:** GitHub Actions, CI/CD, Jenkins, Prometheus, Grafana, ELK Stack, SonarQube

**Data Engineering & Analytics:** Apache Spark, Databricks, Airflow, Trino, Kafka, Kinesis, ETL/ELT, Data Warehousing, SQL.

**Programming Languages:** Python, Rust, SQL, NodeJS, JavaScript, Bash Scripting, CSS.

## Professional Experience

**Senior Platform Engineer** | *Royal Bank of Canada, Toronto*

June 2025 – Present

- **Architected a configuration-driven Data Platform-as-a-Service (dPaaS)** that empowers data engineers to define ETL workflows via YAML. Built connectors for a diverse ecosystem (MongoDB, Postgres, Elasticsearch, Kafka, S3, Kinesis), significantly reducing boilerplate code and accelerating data pipeline delivery.
- **Led a strategic modernization initiative** to migrate legacy applications to GitHub Actions, establishing robust CI/CD pipelines with integrated unit testing, security scanning (Snyk), and code quality gates (SonarQube). Deployed valid artifacts to OpenShift and Azure across **multiple regions to ensure High Availability (HA)** and robust Disaster Recovery protocols
- **Engineered a highly resilient, self-healing application** for distributed Spark workloads. Implemented heuristic-based automation to detect and **autonomously resolve** common failures (OOM, ConfigErrors), reducing operational toil and ensuring platform reliability.

**Senior Cloud Data Engineer** | *Royal Bank of Canada, Toronto*

June 2022 – June 2025

- **Directed the transformation of Databricks into a self-service platform**, enabling 50+ cross-functional partners (Business Insights, Data Science) to deploy models and access curated datasets, fostering a culture of data autonomy.
- **Optimized high-throughput data pipelines** using Apache Airflow and Databricks, aggressively reducing data availability latency from **24 hours to 20 minutes** to support near real-time decision-making apps.
- **Developed and owned a Python-based CI/CD framework** for PowerBI artifacts, automating validation and environment parameterization to ensure seamless releases across Dev, UAT, and Production.
- **Implemented Infrastructure-as-Code (IaC) using Terraform**, reducing environment provisioning time by **80%** for Azure SQL, Data Lake Gen2, and Databricks workspaces.
- **Engineered a real-time embedded visualization application** with Row Level Security (RLS), slashing the report creation cycle **from 4 weeks to 4 hours** while adhering to data governance standards, currently being used by 20 clients.

**Software Developer Co-op** | *Royal bank of Canada, Toronto*

Jan 2020 – Aug 2020

- **Deployed and scaled an Elasticsearch cluster** on OpenShift, configuring high-availability (HA) and fault-tolerance mechanisms to support large-scale logging and monitoring.
- **Built containerized data ingestion pipelines** using Python and Docker, automating the extraction of data from APIs and SharePoint into Elasticsearch to establish a "Single Source of Truth" for operational metrics.

## Project Experience and Industry Involvement

- **Rust Redis Server (2025):** Engineered a high-performance Redis-compatible server using **Rust and gRPC**, supporting persistent storage and concurrent client handling.
- **Spot2YT Music Sync (2024):** Developed a Python automation tool to synchronize Spotify libraries with YouTube Music using OAuth2 and web scraping, handling API rate limits and data validation.
- **P0 Language Extension (2022):** Extended the P0 compiler using **Python and WASM**, implementing formal verification features (preconditions, postconditions, invariants) to enhance code robustness.

## Education

**B.A.Sc. of Computer Science Co-op, McMaster University**

Sept 2017 - Apr 2022